

Machine Learning and Data Mining I: Lecture 4

Linear Classification: Perceptron

Morgan McCarty

10 July 2023

1 Classification

Predict a discrete label rather than a continuous value.

1.1 Artificial Neural Network

- Also known as: Connectionist Model, Adaptive Network
- Inspired by human brain
 - Composed of billions of neurons
 - Fast processing of complex tasks
 - Parallel computations
- Started with Autonomous Land Vehicle (ALVINN) in 1990s.
- Deep Learning:
 - ANNs and variants
 - State of the art for speech recognition
 - Convolutional Neural Networks (CNNs) for image recognition (pattern recognition)
- Tasks:
 - Continuous input features
 - Form of tangent function is unknown
- Drawbacks:
 - Very hard to understand/complex
 - Mostly treated as black box

1.2 Perceptron

- Based on structure of neuron: connect to others using synapses, synapses have varying strength (exciting + or inhibiting −)

1.3 Binary Linear Classification

- Linearly separable Data
- Boundary is linear function of input features (line, plane, hyperplane, etc)
- Input: $x \in \mathbb{R}^m$
- Output: $y \in \{-1, 1\}$

1.4 Perceptron Classification Rule

$$w_1x_1 + w_2x_2 + \cdots + w_mx_m \geq \theta \rightarrow y = 1$$

else: $y = -1$

- w_i is the weight of the i th feature
- x_i is the i th feature
- θ is the threshold (bias)

1.5 Learning Rate

- Learn the $m + 1$ parameters: $w_1, w_2, \dots, w_m, \theta$
- Absorb the threshold into the weight vector: $w_0 = -\theta$ (an extra feature)

$$f(x) = w_0x_0 + w_1x_1 + w_2x_2 + \cdots + w_mx_m$$

where $x_0 = 1$
and $w_0 = -\theta$

1.6 Algorithm

- Data: Training Data = $\{x_i, y_i : \forall i = 1, 2, \dots, n\}$
- Learning Rate: η
- Result: Coefficients of the hyperplane $(w_0, w_1, w_2, \dots, w_m)$
- Algorithm:

- Initialize $w_0, w_1, w_2, \dots, w_m$ to 0
- Repeat until convergence:
 - * For each (x_i, y_i) in training data:
 - Compute $\hat{y}_i = w^T x_i$
 - If $\hat{y}_i y_i \leq 0$ then update w :

$$w = w + \eta y_i x_i$$

- Repeat until $\hat{y}_i y_i > 0$ for all (x_i, y_i)

1.7 Perceptron Learning Example

Learning the boolean *AND* function:

- We have a data table:

x_0	x_1	x_2	$AND(x_1, x_2)$
1	0	0	0
1	0	1	0
1	1	0	0
1	1	1	1

- We want to learn the weights w_0, w_1, w_2 such that $w_0 + w_1x_1 + w_2x_2 \geq 0$ if $x_1 = 1$ and $x_2 = 1$.
- Initialize w_0, w_1, w_2 to 0.
- Set η to 1

– For each (x_i, y_i) in training data:

* Compute $\hat{y}_i = w^T x_i$

* If $\hat{y}_i y_i \leq 0$ then update w :

$$w = w + \eta y_i x_i$$

* Repeat until $\hat{y}_i y_i > 0$ for all (x_i, y_i)

- Let's go through the algorithm step by step for one epoch:

- The first data point is $(1, 0, 0, 0)$

– $\hat{y}_1 = w_0 + w_1x_1 + w_2x_2 = 0$

– $\hat{y}_1 y_1 = 0 \leq 0$

– $w = w + \eta y_i x_i = (0, 0, 0) + 1(0, 0, 0) = (0, 0, 0)$

- The second data point is $(1, 0, 1, 0)$

– $\hat{y}_2 = w_0 + w_1x_1 + w_2x_2 = 0$

– $\hat{y}_2 y_2 = 0 \leq 0$

– $w = w + \eta y_i x_i = (0, 0, 0) + 1(0, 0, 1) = (0, 0, 1)$

- The third data point is $(1, 1, 0, 0)$

– $\hat{y}_3 = w_0 + w_1x_1 + w_2x_2 = 1$

– $\hat{y}_3 y_3 = 0 \leq 0$

– $w = w + \eta y_i x_i = (0, 0, 1) + 1(0, 1, 0) = (0, 1, 1)$

- The fourth data point is $(1, 1, 1, 1)$

– $\hat{y}_4 = w_0 + w_1x_1 + w_2x_2 = 2$

– $\hat{y}_4 y_4 = 1 > 0$

– We do not update w .

- We have completed one epoch (one iteration through the data).

- The current weights are $w = (0, 1, 1)$, $y_w(x) = 0 + 1x_1 + 1x_2 = x_1 + x_2$.
- As we still had errors for the first three data points, we will continue to iterate (we have not converged).
- When we finally have no errors, we have converged.
- After convergence, we have $w = (-4, 3, 2)$, $y_w(x) = -4 + 3x_1 + 2x_2$.
- The exact values of the final weight vector depend on the initial values as well as the learning rate.
- This would not work for the *XOR* function, as it is not linearly separable.

1.8 Questions?

- Will the algorithm always converge for linearly separable data?
 - For two-class linearly separable datasets there is a separating hyperplane w_{opt} .
 - Will the algorithm converge over finite iterations?
- What loss function does the algorithm minimize?
- We assume that the training data has bounded Euclidean norm. (i.e. $\|x_i\| \leq R$ for all i)
- The optimal classifier perfectly separates the data (i.e. $y_i w_{opt}^T x_i > 0$ for all i)

1.9 Convergence

- The algorithm learns the boundary in a finite number of steps based on the number of mistakes it makes starting from w_0 .
- The mistake bound is:

$$k \leq \frac{R^2}{\gamma} \|w_{opt}\|^2 \quad (1)$$

- However most real-world data is not linearly separable.

1.10 Loss to Optimize

- The loss is 0 when the label and prediction agree

$$y_i f_w(x_i) > 0 \quad (2)$$

- The loss is ≥ 0 when the label and prediction values do not agree (i.e. have opposite signs)

$$y_i f_w(x_i) \leq 0 \quad (3)$$

- Therefore the loss function is:

$$L_p(x, y, w) = \max(0, -y w^T x) \quad (4)$$

- Non-zero for misclassified points

1.11 Gradient of Loss Function

- The gradient of the loss function is:

$$L_p = \sum_{i=1}^N \max(0, -y_i w^T x_i)$$

$$L_p = -\sum_{i \in \mu_w} y_i w^T x_i$$

$$\nabla L_p = -\sum_{i \in \mu_w} y_i x_i$$

1.12 Stochastic Gradient Descent

- Batch Gradient Descent

Processing the entire training set for each epoch can be prohibitive for large datasets

- Stochastic Version

Calculate gradient based on individual training instances

1.13 Linearly Inseparable Data

- Can we apply the perceptron to linearly inseparable data?

- Change in criteria:

– Loss: find separator that makes no mistakes \rightarrow find separator that makes least misclassifications (“best fit”)

- Squared Loss:

– Between perceptron output and class labels

– Output: Un-thresholded

– $o(x) = w^T x$

– $E(w) = \frac{1}{2} \sum_{i=1}^N (y_i - o_i)^2$

– $y_i \in \{0, 1\}$

– $o_i \in \mathbb{R}$

– Squared loss increases with values on correct side of boundary